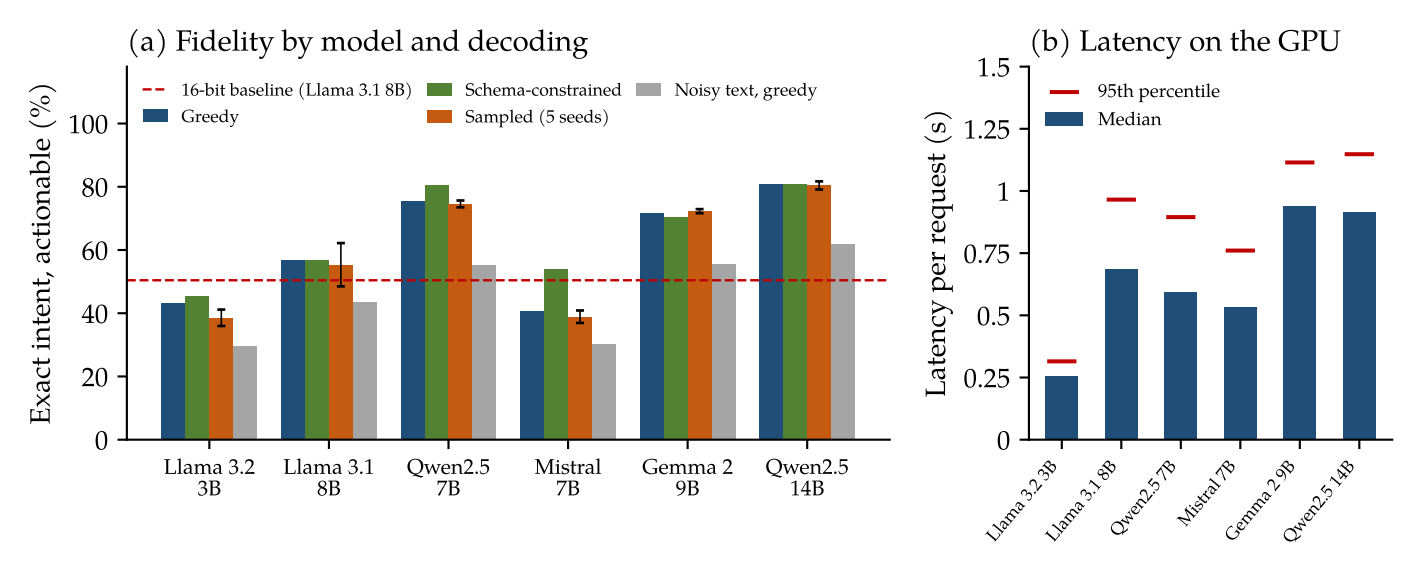
(a) Fidelity by model and decoding
16-bit baseline (Llama 3.1 8B)
Greedy
Schema-constrained
Sampled (5 seeds)
Noisy text, greedy
Exact intent, actionable (%)
0
20
40
60
80
100
Llama 3.2 3B
Llama 3.1 8B
Qwen2.5 7B
Mistral 7B
Gemma 2 9B
Qwen2.5 14B
(b) Latency on the GPU
95th percentile
Median
Latency per request (s)
0
0.25
0.5
0.75
1
1.25
1.5
Llama 3.2 3B
Llama 3.1 8B
Qwen2.5 7B
Mistral 7B
Gemma 2 9B
Qwen2.5 14B